

Supplement A — Thematic Structure

category_diagnostic_profiles.csv

34 rows x 10 columns.

Columns: cluster_id, label, domain, size, alternative_model_recovery, resampling_stability, heldout_replication, geometry_separability, heldout_rep_strict, heldout_rep_direction

cluster_cards.csv

34 rows x 11 columns.

Columns: handle, cluster, size, share_of_clustered, keywords, n_enriched, external_label, exemplar_1_sha1, exemplar_2_sha1, exemplar_3_sha1, exemplars_withheld

partition_stats.csv

1 rows x 6 columns.

Columns: min_cluster_size, n_clusters, noise_pct, clustered_sentences, plateau_lo, plateau_hi

min_cluster_size	n_clusters	noise_pct	clustered_sentences	plateau_lo	plateau_hi
35	34	32.8	5130	30	35

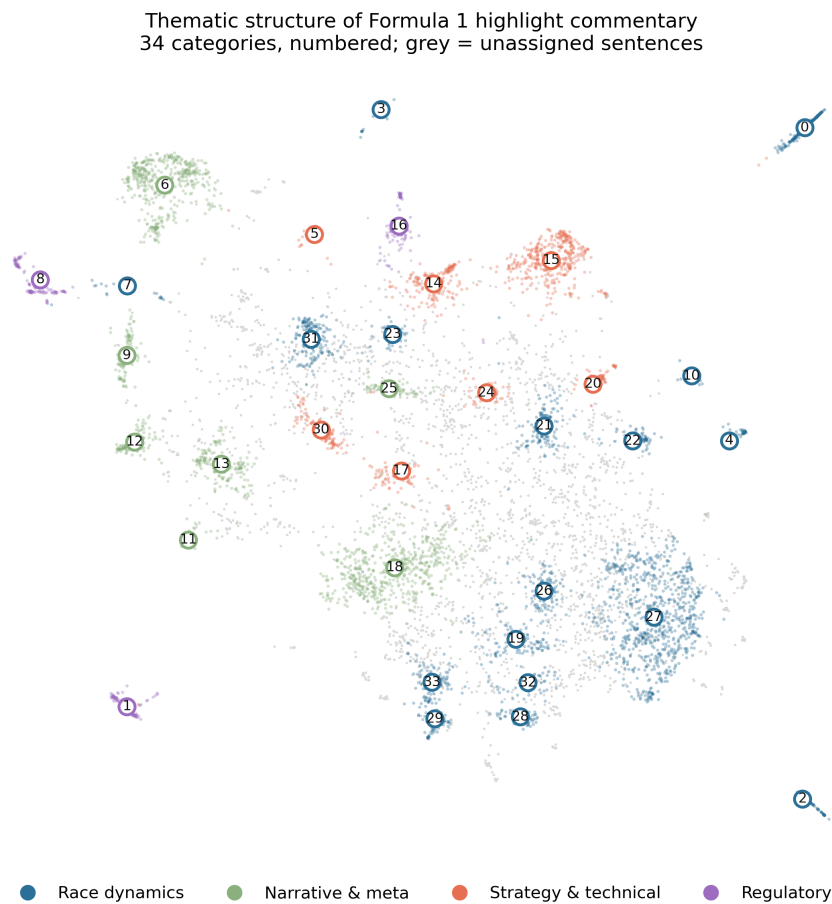
size_concentration.csv

1 rows x 8 columns.

Columns: clustered_sentences, top1_pct, top5_pct, top10_pct, largest, smallest, median, gini

clustered_sentences	top1_pct	top5_pct	top10_pct	largest	smallest	median	gini
5130	15.4	47.1	62.6	792	41	112.5	0.456

Figures



Thematic structure of Formula 1 highlight commentary. Descriptive two-dimensional projection; clustering was performed in the five-component reduced space, not in this plane.